

Top 80 SQL interview questions

1. What is SQL?

- **Answer:** SQL (Structured Query Language) is a standard programming language used for managing and manipulating relational databases.
-

2. What is a primary key?

- **Answer:** A primary key is a field (or combination of fields) that uniquely identifies each record in a table. It cannot contain **NULL** values and must have unique entries.
-

3. What is a foreign key?

- **Answer:** A foreign key is a field (or combination of fields) in one table that refers to the primary key in another table, establishing a relationship between the two tables.
-

4. What are constraints in SQL?

- **Answer:** Constraints are rules applied to table columns to enforce data integrity. Common constraints include:
 - **PRIMARY KEY:** Uniquely identifies each record.
 - **FOREIGN KEY:** Enforces referential integrity.
 - **UNIQUE:** Ensures all values in a column are distinct.
 - **CHECK:** Ensures that values in a column satisfy a specific condition.
 - **NOT NULL:** Ensures a column cannot have **NULL** values.
 - **DEFAULT:** Specifies a default value for a column.
-

5. Write a query to retrieve all records from a table named **employees.**

```
SELECT * FROM employees;
```

- **Answer:** This query selects and displays all columns and rows from the **employees** table.

6. What is the difference between `DELETE` and `TRUNCATE`?

- **Answer:**
 - `DELETE`: Removes rows from a table based on a condition. It can be rolled back (transaction-safe) and triggers can be invoked.
 - `TRUNCATE`: Removes all rows from a table, resetting the identity column. It is faster but cannot be rolled back and does not invoke triggers.
-

7. How do you find the maximum salary from an `employees` table?

```
SELECT MAX(salary) FROM employees;
```

- **Answer:** This query returns the highest salary from the `employees` table.
-

8. Write a query to fetch the second-highest salary from the `employees` table.

```
SELECT MAX(salary) FROM employees  
WHERE salary < (SELECT MAX(salary) FROM employees);
```

- **Answer:** The subquery finds the maximum salary, and the outer query finds the highest salary that is less than that value (i.e., the second-highest salary).
-

9. What is a `JOIN`? Explain its types.

- **Answer:** A `JOIN` clause is used to combine rows from two or more tables based on a related column. Types of joins:
 - `INNER JOIN`: Returns rows with matching values in both tables.
 - `LEFT JOIN`: Returns all rows from the left table and matching rows from the right.
 - `RIGHT JOIN`: Returns all rows from the right table and matching rows from the left.

- o FULL JOIN: Returns rows when there is a match in either table.
 - o CROSS JOIN: Returns the Cartesian product of both tables.
-

10. Write a query to fetch employee names and department names using JOIN.

```
SELECT e.name, d.department_name
FROM employees e
JOIN departments d
ON e.department_id = d.id;
```

- **Answer:** This query joins the `employees` table with the `departments` table based on the `department_id`, displaying employee names and their corresponding department names.
-

11. What is a GROUP BY clause in SQL?

- **Answer:** The `GROUP BY` clause groups rows with the same values into summary rows. It is commonly used with aggregate functions like `COUNT()`, `SUM()`, `AVG()`, etc.
-

12. Write a query to count employees in each department.

```
SELECT department_id, COUNT(*)
FROM employees
GROUP BY department_id;
```

- **Answer:** This query groups employees by `department_id` and counts the number of employees in each department.
-

13. What is the difference between WHERE and HAVING clauses?

- **Answer:**

- **WHERE**: Filters rows before grouping (applies to individual rows).
 - **HAVING**: Filters groups after the GROUP BY clause (applies to aggregate functions).
-

14. Write a query to fetch departments with more than 5 employees.

```
SELECT department_id, COUNT(*)  
FROM employees  
GROUP BY department_id  
HAVING COUNT(*) > 5;
```

- **Answer:** The query counts employees in each department and returns departments with more than 5 employees.
-

15. Explain UNION and UNION ALL.

- **Answer:**
 - **UNION**: Combines results of two or more SELECT statements and removes duplicates.
 - **UNION ALL**: Combines results and keeps all duplicates.
-

16. What is a subquery in SQL?

- **Answer:** A subquery is a query nested within another query. It is used to retrieve data that will be passed into the outer query.
-

17. Write a query to find all employees whose salary is greater than the average salary.

```
SELECT *  
FROM employees  
WHERE salary > (SELECT AVG(salary) FROM employees);
```

- **Answer:** This query selects all employees with a salary higher than the average salary of all employees.
-

18. What is the difference between INNER JOIN and OUTER JOIN?

- **Answer:**
 - **INNER JOIN:** Returns rows with matching values in both tables.
 - **OUTER JOIN (Left/Right/Full):** Returns matching rows plus non-matching rows from one or both tables.
-

19. Write a query to fetch the current date in SQL.

```
SELECT CURRENT_DATE;
```

- **Answer:** This query retrieves the current date from the database.
-

20. What is indexing in SQL?

- **Answer:** Indexing improves the speed of data retrieval by creating a data structure (index) on one or more columns of a table.
-

21. What is normalization? Explain its types (1NF, 2NF, 3NF, BCNF).

- **Answer:** Normalization is the process of organizing data to reduce redundancy and improve data integrity. Forms:
 - **1NF:** Eliminate duplicate columns and create tables for related data.
 - **2NF:** Remove partial dependencies (columns depend on a part of a composite key).
 - **3NF:** Remove transitive dependencies (non-key columns depend on other non-key columns).
 - **BCNF:** A stricter version of 3NF where every determinant must be a candidate key.
-

22. What is denormalization?

- **Answer:** Denormalization is the process of combining normalized tables to improve performance at the cost of introducing redundancy.
-

23. Write a query to add a new column `email` to the `employees` table.

```
ALTER TABLE employees ADD COLUMN email VARCHAR(255);
```

- **Answer:** This query adds a new `email` column to the `employees` table.
-

24. What is a stored procedure in SQL?

- **Answer:** A stored procedure is a set of SQL statements that can be stored in the database and executed as a program to perform a specific task.
-

25. Write a basic stored procedure to fetch all employees.

```
CREATE PROCEDURE GetAllEmployees()
BEGIN
    SELECT * FROM employees;
END;
```

- **Answer:** This procedure retrieves all records from the `employees` table when executed.
-

26. What are triggers in SQL?

- **Answer:** Triggers are special procedures that are automatically executed (or "triggered") in response to certain events (INSERT, UPDATE, DELETE) on a table.
-

27. Write a query to create a trigger that logs any delete action on the employees table.

```
CREATE TRIGGER log_delete
AFTER DELETE ON employees
FOR EACH ROW
BEGIN
    INSERT INTO log_table(action, emp_id, log_time)
    VALUES('DELETE', OLD.id, NOW());
END;
```

- **Answer:** This trigger logs the deletion of any employee by recording the action and employee ID in the `log_table`.
-

28. What is a VIEW in SQL?

- **Answer:** A `VIEW` is a virtual table based on the result set of an SQL query. It does not store the data itself but provides a way to simplify complex queries.
-

29. Write a query to create a view for employees with salary greater than 50,000.

```
CREATE VIEW HighSalaryEmployees AS
SELECT * FROM employees WHERE salary > 50000;
```

30. What is the difference between VIEW and TABLE?

- **Answer:** A `TABLE` stores data physically, while a `VIEW` is a virtual representation that dynamically pulls data from one or more tables without storing it.
-

31. What is an aggregate function? Provide examples.

- **Answer:** Aggregate functions perform calculations on a set of values and return a single value. Examples include:
 - COUNT(): Counts the number of rows.
 - SUM(): Sums up a numeric column.
 - AVG(): Calculates the average of a numeric column.
 - MAX(): Returns the maximum value.
 - MIN(): Returns the minimum value.
-

32. Write a query to calculate the total salary for each department.

```
SELECT department_id, SUM(salary)
FROM employees
GROUP BY department_id;
```

- **Answer:** This query sums the salaries for each department, grouping by department_id.
-

33. Explain the DISTINCT keyword in SQL.

- **Answer:** The DISTINCT keyword is used to return unique values from a column, eliminating duplicate entries from the result set.
-

34. Write a query to find distinct job titles from the employees table.

```
SELECT DISTINCT job_title FROM employees;
```

- **Answer:** This query retrieves unique job titles from the employees table.
-

35. What are the ACID properties in SQL?

- **Answer:** ACID properties ensure reliable processing of database transactions:

- **Atomicity:** Ensures that all parts of a transaction are completed successfully or none at all.
 - **Consistency:** Ensures the database remains in a valid state before and after the transaction.
 - **Isolation:** Ensures transactions do not affect each other's execution.
 - **Durability:** Ensures that once a transaction is committed, it remains permanent, even in the event of a failure.
-

36. What is a transaction in SQL?

- **Answer:** A transaction is a sequence of one or more SQL operations treated as a single unit of work, ensuring data integrity.
-

37. Explain COMMIT, ROLLBACK, and SAVEPOINT.

- **Answer:**
 - **COMMIT:** Saves all changes made during the current transaction.
 - **ROLLBACK:** Undoes changes made during the current transaction, restoring the database to its previous state.
 - **SAVEPOINT:** Sets a point within a transaction to which you can later roll back.
-

38. Write a query to start a transaction, update a record, and commit it.

```
START TRANSACTION;  
UPDATE employees SET salary = 60000 WHERE id = 1;  
COMMIT;
```

- **Answer:** This sequence starts a transaction, updates an employee's salary, and commits the change.
-

39. What is a CASE statement in SQL?

- **Answer:** A **CASE** statement is used to perform conditional logic in SQL queries, allowing different outputs based on specified conditions.

40. Write a query using CASE to categorize employees by salary.

```
SELECT name,
       CASE
           WHEN salary > 50000 THEN 'High'
           WHEN salary BETWEEN 30000 AND 50000 THEN 'Medium'
           ELSE 'Low'
       END AS salary_category
FROM employees;
```

- **Answer:** This query categorizes employees based on their salary levels.
-

41. Explain NULL values in SQL.

- **Answer:** NULL represents the absence of a value in a database. It is not equivalent to zero or an empty string and is treated differently in comparisons.
-

42. Write a query to fetch records where email is NULL.

```
SELECT * FROM employees WHERE email IS NULL;
```

- **Answer:** This query retrieves all employees whose email address is not provided (i.e., is NULL).
-

43. What is the COALESCE function in SQL?

- **Answer:** The COALESCE function returns the first non-NULL value in a list of expressions.
-

44. Write a query using COALESCE to handle NULL values in a column.